


# llama3 16-rank two-tier A/B: endpoint vs INC per rendering

scale-out tree16 (100 Gbps, lossy) + per-node single-switch scale-up (3600 Gbps, lossless PFC)



each pair = controlled within-rendering A/B; plain → SP is a workload change (2× TP collectives, different baseline and compute) — the -4.65 → +6.01 % flip at C3 is not like-for-like (placeholder compute model)