


llama3 16-rank two-tier A/B: endpoint vs INC per rendering
scale-out tree16 (100 Gbps, lossy) + per-node single-switch scale-up (3600 Gbps, lossless PFC)



each pair = controlled within-rendering A/B; plain $\rightarrow$ SP is a workload change (2 $\times$ TP collectives, different baseline and compute) — the 2.40 $\rightarrow$ 14.15 % jump at C3 is not like-for-like