


llama3 16-rank two-tier A/B: endpoint vs INC per rendering  
scale-out tree16 (100 Gbps, lossy) + per-node single-switch scale-up (3600 Gbps, lossless PFC)



each pair = controlled within-rendering A/B; plain  $\rightarrow$  SP is a workload change (2 $\times$  TP collectives, different baseline and compute) — the -0.52  $\rightarrow$  +10.39 % swing at C3 is not like-for-like (placeholder compute model)